


wasocaml: compiling OCaml to Wasm

Pierre Chambart <pierre.chambart@ocamlpro.com>¹

Léo Andrès <l@ndrs.fr>^{1, 2}

January 10, 2023

1. OCamlPro
2. Université Paris-Saclay, CNRS, ENS Paris-Saclay, Inria, Laboratoire Méthodes Formelles



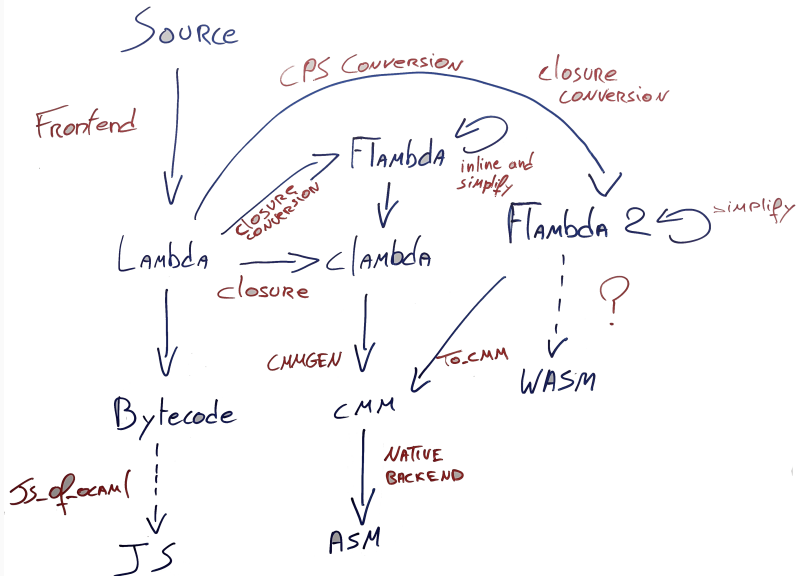
- we're programming languages consultants
- specialized in OCaml, Rust and formal methods
- we develop flambda/flambda2 (ocaml optimizer IR)

OCaml and Wasm

- `js_of_ocaml`: OCaml bytecode to JavaScript
- `wasicaml`: OCaml bytecode to Wasm

- functions can be: mutually recursive, nested, polymorphically recursive, first class, partially applied
- algebraic datatypes, GADTs, polymorphic variants, objects
- exceptions
- modules, functors, first class modules
- allow low-level manipulations through the `Obj` module
- GC with good performances (minor/major heaps)

OCaml compilation chain



IR choice

- we chose to go from Flambda to Wasm (we'll use Flambda2 in the future)
- why? we don't want to rewrite the compiler (OCaml typechecking is hard)

Values representation

- uniform representation
- 1 bit to distinguish between scalars and pointers to heap-allocated blocks
- a given type can have values of both kinds:
 - `Obj.is_int (Obj.repr []) = true`
 - `Obj.is_block (Obj.repr [1; 2]) = true`

Flambda...

- ANF
- explicit closures
- high-level: works on abstract values and not directly on the actual memory layout (this is done by Cmm)

Block compilation

Two strategies:

- before: `struct { i16 size; i8 tag; refeq data }`
- didn't work because of too long subtyping chains
- now: `refeq array (tag followed by data)` (size is not needed)
- we kept both and can switch easily
- we could use more precise types with the `struct` strategy

...to Wasm

```
(type $Gen_block (struct (field (mut i8)) (field (mut i16))))  
(type $Block_0  
  (struct_subtype (field (mut i8)) (field (mut i16))  
    $Gen_block))  
(type $Block_1  
  (struct_subtype (field (mut i8)) (field (mut i16)) (field  
    (mut (ref eq))))  
  $Gen_block))  
(type $Block_2  
  (struct_subtype (field (mut i8)) (field (mut i16)) (field  
    (mut (ref eq))))  
  (field (mut (ref eq))) $Block_1))
```

...to Wasm

```
(type $Float (struct (field (mut f64))))
(type $Int32 (struct (field (mut i32))))
(type $Int64 (struct (field (mut i64))))
(type $String (array (mut i8)))
(type $Array (array (mut (ref eq))))
(type $FloatArray (array (mut f64)))
(rec
  (type $Func_1
    (func (param (ref eq)) (param (ref $Env)) (result (ref
      eq)))))
  (type $Env (struct (field (mut i8)) (field (mut (ref
    $Func_1)))))
(type $Func_2
  (func (param (ref eq)) (param (ref eq)) (param (ref $Env))
    (result (ref eq))))
(type $Gen_block (array (mut (ref eq))))
```

...to Wasm

```
(type $Gen_closure_3
  (struct_subtype (field (mut i8)) (field (mut (ref
    $Func_1)))
    (field (mut (ref $Func_3))) $Env))
(type $Closure_3_1
  (struct_subtype (field (mut i8)) (field (mut (ref
    $Func_1)))
    (field (mut (ref $Func_3))) (field (mut (ref eq)))
      $Gen_closure_3))
```

Why do we need i31

- two cases: `int` and *small scalars*
- GADTs are complicated:

```
type _ tag =  
  | I : int tag  
  | F : float tag  
  | Box : (int * int) tag  
  
let f : type t. t -> t tag -> int =  
  fun v tag ->  
    match tag with  
    | I -> v  
    | F -> int_of_float v  
    | Box -> fst v + snd v
```

Difficulties

- FFI probably won't work with emscripten (in `js_of_ocaml` one needs to write bindings by hand)
- partial application
- too long subtyping chain
- objects (possible but not done yet)
- exceptions: we can't use wasm exceptions as OCaml allows dynamic exception creation (through functors), so we only have one wasm exception and then we use our own identifiers for everything else

Things that went well

- GC proposal well designed
- i31 really helped
- people were very reactive (e.g. closed world assumption by binaryen)
- OCaml compilation is fast, so having one Wasm module per OCaml module makes separate compilation of the two match quite well (but we have to export a lot of globals)

Benchmarks

compiler	fib 36	fib 37	ocamlsh test
ocaml native	0.14	0.21	1.83
ocaml bytecode	0.72	1.21	3.21
js_of_ocaml	0.29	0.40	4.98
wasicaml	0.55	0.85	3.65
ocamlrun wasm	4.09	6.50	8.85
wasocaml	0.31	0.42	3.54

- fib: `js_of_ocaml` is OK
- ocamlsh test: `js_of_ocaml` is not good enough, but we are
- `js_of_ocaml` is sometimes up to x40 times slower than native in real examples

- OCaml is now multicore
- `js_of_ocaml` is going to use CPS+trampoline
- it's too costly for us
- need for stack switching (effects handlers)

Future

- `let` will probably be useful
- we won't need `wasm` closures
- most of our instructions are casts, `binaryen` removes some, `flambda2` will remove much more; for others we'll have to propagate types further in the compiler or generate better code
- we probably won't need `rtt` but they may lead to better performances (need to think more about it)
- experiment with whole-program linking (`binaryen-merge` would be useful), we won't use `ld` from `llvm` as we don't have `.o` files