

# Stack Switching

Exploring the Design Space

Andreas Rossberg 2022/10/27

# Typed Continuations Proposal

Arjun Guha, Daniel Hillerström, Daan Leijen, Sam Lindley, Luna Phipps-Costin, Matija Pretnar, KC Sivaramakrishnan

Still very actively being worked on!

... Daniel & Luna implementing in [Wasmtime](#), first prototype almost complete

... Daniel's talk on Tuesday

Design has been fairly [stable](#), a collection of non-trivial examples have been written

We have focussed on creating the basis for actual [evaluation](#) and tangible data

... which takes time, e.g., needed to implement prerequisites like funcrefs first

Have been rather bad at PR and giving updates :)

**Convergence so far**

# Delimited Continuations

all current suggestions implement a form of *delimited continuations*

- ... *hierarchical* stacks

- ... *structured*, clean interop with exceptions and typing

- ... more *powerful* than undelimited continuations

- ... in *practice*, typically need at least an outermost delimiter

  - e.g., even around threads, to guard event loop in the browser, REPL in an interpreter, ...

- ... *proven* in Wasmtime fibers, Wasm/k

- ... what many other communities have converged on as well



# Interlude: Terminology

many competing or overlapping uses of terminology

... e.g., Java 19 *continuation*  $\approx$  Wasmtime *fiber*

sometimes reflect **semantic** vs **implementation** point of views

trying to stick to:

**computation**: sequential execution of a subprogram that can be suspended

**continuation**: remainder of a computation when suspended

**stack**: system resource to run a computation on

note that **fiber** is often used interchangeably with all three, so avoiding it for this presentation

# Core Primitives

three main primitives

**create** : (ref \$func)  $\rightarrow$  (ref \$suspension)

**resume** : (ref \$suspension)  $\text{arg}^* \rightarrow \text{res}^*$     ;; plus **handler**

**suspend** : (ref \$target)  $\text{arg}^* \rightarrow \text{res}^*$

computations are expressed by an initial function

... this should be a perfectly ordinary function

# Handlers

every **resume** is associated with a **handler**

basically a jump target or **jump table** for suspension **events**

mainly superficial syntactic differences in latest proposals

```
(block $l
  ...
  (resume (event $tag $l))
  ...
)
... ;; event landing pad
```

```
(handler
  ...
  (resume)
  ...
  (event $tag
    ... ;; event landing pad
  )
)
```



**The design space from here on**



# Main Dimensions

typing of values

handler selection

switching mechanisms

# Additional Considerations

memory management

cancellation and clean-up


performance overhead

implementation complexity

capabilities?

# Typing

# Lifetime of a Computation



suspends are dealt with by a **handler** (a.k.a. prompt)

but what are the **types** of  $a^*$  and  $r^*$ ?

... a **spectrum** of options with **increasing precision**, decreasing overhead

... in all relevant use cases, types are determined by **event** (event emit  $\approx$  function call)



# Event Examples

## threads

yield : () → ()  
spawn : (func () ()) → i32  
self : () → i32


## async/await

await : (ref \$fut) → anyref  
async : (func () anyref) → (ref \$fut)

## actors

send : (ref \$actor) anyref → ()  
recv : (ref \$actor) → anyref  
self : () → (ref \$actor)  
new : (func () ()) → (ref \$actor)

# Option 0: no values




no values can be passed along resume/suspend

... e.g., traditional coroutines, Java 19

parameters and results must be funnelled through global mutable state

... undesirable for the abstractions we are interested in; problematic with threading

# Option 1: homogeneous types



all  $a^*$  and all  $r^*$  have the **same type**


... e.g., classic shift/reset or control/prompt in untyped languages; Francis' second proposal

in practice, parameters must **encode a union**

... difficult and inefficient with Wasm value types; often implies boxing; non-nullable refs



# Option 1b: mostly homogeneous types



all  $a^*$  and all  $r_i^*$  have the **same type**, but  $r^*$  is typed **separately**

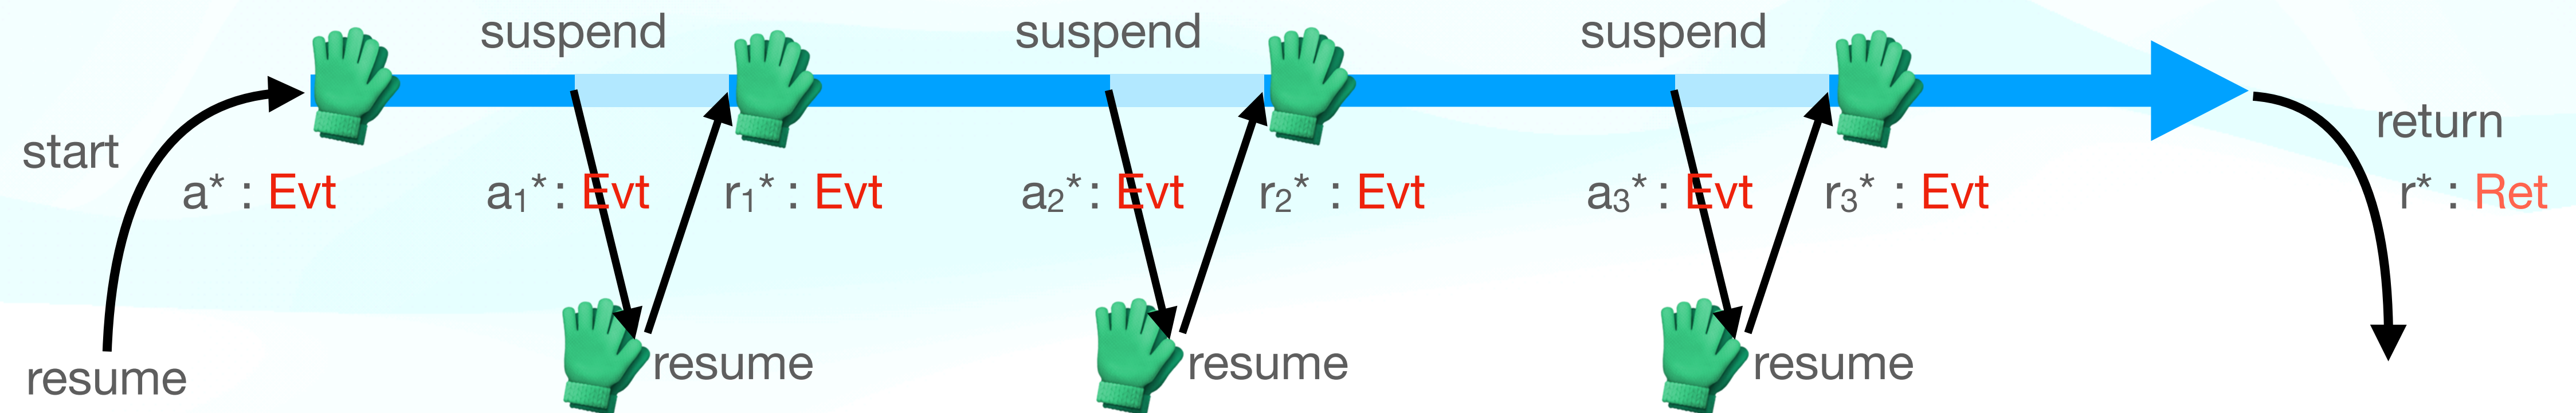
... e.g., Wasmtime fibers

in practice, parameters must **still encode a union**

... difficult and inefficient with Wasm value types; often implies boxing; non-nullable refs



# Option 2: dynamic types



all values have **same/no static type** but can have **different dynamic type**

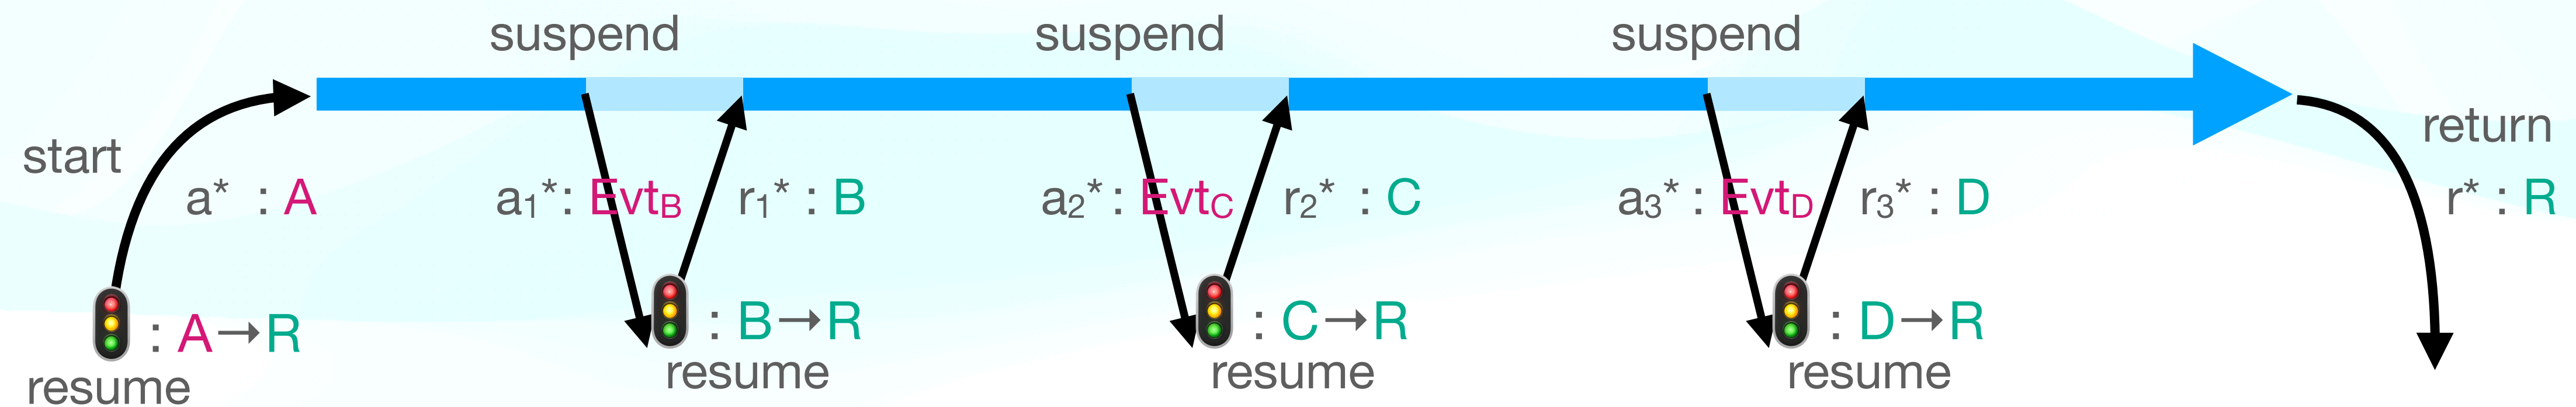
... Francis' third proposal; no precedent

requires dynamic **check/dispatch/handler everywhere** (start, suspend, resume)

... runtime cost; substantial complexity (e.g., every function needs an extra entry point)

... not a natural fit for function-like nature of events; entry handler for computation is unorthodox

# Option 3: heterogeneous types



all values can have **different static type**, each  $a_i^*$  type **implies** respective  $r_i^*$  type

... e.g., typed continuations proposal; effect handlers

requires refining fibers to explicit **continuations** to distinguish resumption types


... check needed to ensure they are used only once

... but no allocation needed for them (this was a valid concern!)

# Handler Selection



# Topology of Stacks and Handlers




**suspend** transfers control to an **active handler** (arrows correspond to handlers, active when connected)

**resume** transfers control to a **continuation** (an unconnected stack and its children)



# Topology of Stacks and Handlers



(1) suspension [target](#) is a [handler](#), not a stack


(2) suspension is only legal to handlers currently in (dynamic) scope

... handler lookup is [dynamic scoping](#), in **all** proposals! (visible in both [semantics](#) and [implementation](#))

... even if the “name” of the handler is known, its relative “binding” point is not


... different implementation strategies for dynamic scoping, but all have linear cost in this setting

# Selecting handlers



spectrum of options for identifying a handler

# Option 1: innermost handler




suspends unconditionally to [innermost](#) handler (a.k.a. [single prompt](#))

... classic [shift/reset](#) or [control/prompt](#); [Wasm/k](#); [Wasmtime fibers](#) (?)

does not compose well for multiple different control abstractions

# Option 2: by handler reference




suspends to the [referenced](#) handler – if in scope

... precedent?

need to funnel reference everywhere



# Option 2b: by stack reference




suspends to handler **currently associated** with respective stack

... task/fibers proposals; precedent?

have to funnel ref everywhere; additional failure cases, because stacks and handlers are not in a 1:1 relation

# Option 3a: by handler label




suspends to the [innermost](#) handler with respective label (a.k.a. [multi prompt](#)); labels defined separately

... [Java 19](#) runtime; modern [Scheme/Lisp](#) libraries; [libmprompt](#)

label can be static or first-class; subsumes by-reference if labels can be first-class and generative

this is a highly relevant point in the design space, worth exploring!

# Option 3b: by handler label set




suspends to the **innermost** handler that includes respective label

... **effect handlers**; Koka; typed continuations proposal

improves composability and typing precision over single label

# Option 3c: by label predicate




suspends to innermost handler that [accepts](#) label

... [Ocaml](#) runtime

requires closures in general; probably less optimisable



# Option 1+2: by handler reference & label set



suspends to [referenced](#) handler's branch with [respective label](#)

improves static typing over mere handler reference

still have to funnel ref everywhere; superseded by first-class labels

# Direct Switching

# Direct switch

additional primitive, roughly, a combination of suspend+resume

**switch** : (to : susp  $t_1^*$   $t_2^*$ ) (from : target  $t_3^*$   $t_2^*$ )  $t_1^* \rightarrow t_4^*$

replaces the target computation (current or parent) with another computation

motivation: bypass one hop

switch : suspend/resume  $\approx$  tailcall : return/call



# Direct switch

can be combined with any of the other choices discussed so far

complex, since it involves three stacks at the same time

“undermines” structured concurrency by unconditionally bypassing handler

... should handlers at least be able to decide whether they allow this?

unclear if it actually is a relevant performance win

premature optimisation? need to measure

# Other Considerations

# Memory Management

**fibers** depend on mutation of higher-order state and immediately allow for **heap cycles**

```
(type $f (func (param anyref)))  
(type $s (fiber $f))  
  
(func $f (param $x anyref) (suspend) ...)  
  
(func $mk-cyclic-fiber (result (ref $s))  
  (local $s (ref null $s))  
  (local.set $s (fiber.new $s (ref.func $f)))  
  (resume (local.get $s) (local.get $s))  
  (local.get $s)  
)
```

engines that need to prevent leaks from abandoned fibers **face full GC problem**

... can't clean up after failed or abandoned modules without GC

**continuations** o.t.o.h. can **avoid cycles**, enabling RC

... implementation choice: trade-off between reusing memory and simplifying memory management



# Cancellation and Clean-up

stacks may capture (runtime-internal) resources

to abort a stack, it should be properly **unwound**

... every stack ought to be used **linearly** until it either terminates or throws

simply **deleting** a stack is almost always **wrong**

... retire/release are footguns

**exception mechanism** already provides all the necessary functionality

# Performance

We don't know much yet, mainly guess work

Need [experimental evaluation](#)

... we have been focussing on implementation in [Wasmtime](#) lately

... complementary experiments by [V8](#) team

Should avoid [premature optimisation](#) in the design at this stage

# Capabilities?

*“What is your unit of isolation?”* – Mark Miller

proper capability for **resume** is a **continuation**

proper capability for **suspend** is a **handler reference** or a **first-class tag**

**stacks** or **fibers** are *not* capability-oriented

... too **coarse-grained**: don't distinguish suspend ( $\approx$  receive) from resume ( $\approx$  send)

... too **long-lived**: once given, can be used arbitrarily later, for the whole lifetime



# Use Cases

# Use Cases

	events	handlers	cancellation policy
generators	1, homogeneous	local, nested	unwind
async/await	1+, heterogeneous	varies (event loop, micro events, other future notions)	unwind or other
threads	2+, heterogeneous	global	unwind
actors	3+, heterogeneous	global	none?
control/prompt	1, mostly homogeneous	local, nested	unwind+
call/cc	1, homogeneous	global	?
effect handlers	N, heterogeneous	local, nested	unwind

# Summary

# Next Steps

Much is assumed, little is known!

Need to converge further somehow

Typed continuations proposal takes cues from observable trend in this field

... undelimited continuations → delimited continuations → multi-prompt → effect handlers

(these go by different names in different communities, e.g., symmetric vs asymmetric coroutines, fibers, task scopes, etc)

We have many data points from other languages, but not much for Wasm

Our current focus needs to be [experimentation](#) and [evaluation](#)

... of both [performance](#) and [usability](#)



# Backup

# Representation of Continuations

`fiber` = stack + cpu state + parent fiber + i64 sequence counter

`continuation` = fiber + i64 sequence counter

`resume` : (ref \$continuation) args\*  $\rightarrow$  res\*


- ... check that sequence counter matches up

- ... increment sequence counter in fiber

a continuation can be represented as a `fat pointer`, no allocations required

- ... natural choice now that we have multiple reference hierarchies

# Topology of Stacks and Handlers



(1) resumption [target](#) is a [continuation](#), encapsulating a (disconnected) *chain* of stacks, not a single stack

(2) execution proceeds on the top stack, not the disconnected bottom stack

... `resume(fiber)` and `switch(fiber)` are misleading, as they generally switch to a child of the fiber



# The Problem with by-reference

various control abstractions can be **nested**

- ... generators, control/prompt, some forms of async/await, effect handlers, exception handlers, ...

when suspending, need to get hold of the **current** handler reference for the event at hand

naturally a problem of **dynamic scoping**

- ... which the engine already implements!

- ... but asking for a reference asks producer to **replicate** that mechanism in user space

- ... language with N control effects generally requires N dynamically scoped bindings

# Replicating Dynamic Scoping in User Space

**pure** approach: add N pervasive **parameters** to calling conventions

- ... or one N-tuple, but that trades off more parameters for more allocations

- ... significant call **overhead** that affects *all* functions!  
(unless source language has a static type & effect system to distinguish, e.g., Koka)

**impure** approach: maintain a global **shadow stack** – actually a **forest** of stacks

- ... complex and brittle: have to keep two dynamic scoping mechanisms in sync

- ... global state interacts poorly with other effects like threads and exceptions

- ... unnecessarily duplicates work the engine already does

- ... *did anybody ever try this?*

neither approach is **cheap**, neither is **composable**!

# Selecting handlers by reference

As far as we can tell, this approach is

- ... **complicated** & error-prone
- ... **monolithic** and not composable
- ... more **expensive** than necessary
- ... **unproven**, negating the experience from other language communities
- ... without tangible advantage



**Thank you**